

Asynchronous Effect Handling

Leo Poulson

Master of Science
Computer Science
School of Informatics
University of Edinburgh
2020

Abstract

Features for asynchronous programming are commonplace in the programming languages of today, allowing programmers to issue tasks to run on other threads and wait for the results to come back later. This is particularly useful for programs like web programs, etc...

In this thesis we show how asynchronous programming can be very easily accommodated in a language with existing support for effect handlers. We show how, with a small change to the language implementation, truly asynchronous programming with pre-emptive concurrency is achieved.

Acknowledgements

thanks!

Table of Contents

1	Programming in Frank	1
1.1	Effects and Effect Handling	1
1.2	Concurrency	4
1.2.1	Simple Scheduling	4
1.2.2	Forking New Processes	5
2	Formalisation of Frank	8
3	Pre-emption	14
3.1	Motivation	14
3.2	Interruption with Yields	14
3.3	Relaxing Catches	16
3.4	Interrupting Arbitrary Terms	17
3.5	Interrupting In Practice	19
	Bibliography	20
A	Remaining Formalisms	21

Chapter 1

Programming in Frank

Frank is a functional programming language, designed with the use of algebraic effects at its heart. As such, Frank has an effect type system used to track which effects a computation may use.

Frank also offers very fine-grained control over computations. It clearly distinguishes between computation and value, and offers *multihandlers* to carefully control when computations are evaluated. This combined with effect handling provides a very rich foundation for expressing complex control structures.

In this chapter, we introduce the language, and show why it is so well-suited to our task. We assume some familiarity with typed functional programming, and skip over some more traditional aspects of the language — algebraic data types, pattern matching, etc — so we can spend more time with the novel, interesting parts of the language.

1.1 Effects and Effect Handling

Interfaces and Operations Frank encapsulates effects through *interfaces*, which offer *commands*. For instance, the **State** effect (interface) offers two operations (commands), **get** and **put**. In Frank, this translates to

```
interface State X = get : X
                  | put : X -> Unit
```

The type signatures of the operations mean that **get** is a 0-ary operation which is *resumed* with a value of type **X**, and **put** takes a value of type **X** and is resumed with **unit**. Programs get access to an interface's command by including them in the *ability* of the program. Commands are invoked just as normal functions;

```

xplusplus : {[State Int] Unit}
xplusplus! = put (get! + 1)

```

This familiar program increments the integer in the state by 1.

Handling Operations Traditionl functions in Frank are a specialisation of Frank's handlers; that is to say, functions are handlers that handle no effects. A handler for an interface pattern matches *on the operations* that are invoked, as well as on the *values* that the computation can return. Furthermore, the handler gets access to the *continuation* of the calling function as a first-class value. Consider the handler for **State**;

```

runState : {<State S> X -> S -> X}
runState <get -> k>   s = runState (k s) s
runState <put s -> k> _ = runState (k unit) s
runState x           _ = x

```

The type of **runState** expresses that the first argument is a computation that can perform **State** **S** effects and will eventually return a value of type **X**, whilst the second argument is a value of type **S**. The **State** **S** effect is then *removed* from the first argument.

What happens when we run **runState xplusplus! 0**? When a computation is invoked, it is performed until it results in either a *value* or a *command*. Thus, **runState** will be paused until **xplusplus!** reduces; **runState** is resumed when **xplusplus** is in one of these two forms.

xplusplus instantly invokes **get!**. At this point, control is given to the handler **runState**; both in the sense that **runState** is now being executed by the interpreter, and that **runState** has control over the *continuation* of **xplusplus**, which is a function of type **Int -> [State Int] Unit**. We see that **runState** chooses to resume this continuation with the value of the state at that time.

Top-Level Effects Some effects need to be handled outside of pure Frank, as Frank is not expressive or capable enough on its own. Examples are console I/O, web requests, and ML-style state cells.

These are handled by the interpreter.

Synchronicity and Conversations Observe how the interaction between the effect invoking function and the handler of this effect becomes like a conversation; the caller

asks the handler for a response to an operation, and the caller will then wait, blocking, for a response. This can be characterised as *synchronous* effect handling.

But what if we want to make a request for information, then do something else, then pick up the result later when we need it? This is the canonical example of asynchronous programming. It is not as simple as just invoking our e.g. `getRequest` effect; computation would block once this is invoked, meaning we are stuck waiting for the request to return.

This asynchrony is exactly what we search for in this project.

Multihandlers Recall that in Frank pure functions are just the special case of handlers that handle no effects. Naturally, this notion extends to the n -ary case; we can handle multiple effects from different sources at once. Handlers which handle multiple effects simultaneously are unsurprisingly called *multihandlers*. This lets us write functions such as `pipe` (example due to Convent et al. [2020]);

```

1 interface Send X = send : X -> Unit
2
3 interface Receive X = receive : X
4
5 pipe : {<Send X>Unit -> <Receive X>Y -> [Abort]Y}
6 pipe <send x -> s> <receive -> r> = pipe (s unit) (r x)
7 pipe <_> y = y
8 pipe unit <_> = abort!

```

Line 5 states that `pipe` will handle all instances of the `Send` effect in the first argument, all instances of the `Receive` effect in the second, and might perform `Abort` commands along the way. The matching clauses are also new to the reader; line 6 implements the communication between the two functions. We reinvoke `pipe`, passing the payload `x` of `send` to the continuation of `r`. Lines 7 and 8 make use of the *catch-all* pattern, `<m>`. This will match the invocation of any effect that is handled by that argument, or a value, binding this to `m`. In line 7, the catchall pattern matches either a `send` command or a value; in this case, the receiver has produced a value, so we can return that. In line 8 `<_>` matches either a value or a `receive`; but it must be a `receive` command, as the value case would have been caught above. Hence we have a broken pipe, so the `abort` command is invoked. This can then be caught by another handler, which can implement a recovery strategy.

TODO: Is it worth changing the example to match request on the left earlier?

1.2 Concurrency

Frank is a single-threaded language. It is fortunate, then, that effect handlers give us a malleable way to run multiple program-threads “simultaneously”

TODO: This is poorly written — fix

This is because the invocation of an operation not only offers up the operation’s payload, but also the *continuation* of the calling computation. The handler for this operation is then free to do what it pleases with the continuation. For many effects, such as `getState`, nothing interesting happens to the continuation; in the case of `getState`, it is resumed with the value in state. But these continuations are first-class; they can be resumed, sure, but also stored elsewhere or even thrown away. As such, by handling `Yield` operations, we easily pause and switch between several threads.

1.2.1 Simple Scheduling

We introduce some simple program threads and some scheduling multihandlers, to demonstrate how subtly different handlers generate different scheduling strategies.

```
interface Yield = yield : Unit
```

```
words : {[Console, Yield] Unit}
```

```
words! = print "one "; yield!; print "two "; yield!; print "
  three "; yield!
```

```
numbers : {[Console, Yield] Unit}
```

```
numbers! = print "1 "; yield!; print "2 "; yield!; print "3 ";
  yield!
```

First note the simplicity of the `Yield` interface; we have one operation supported, which looks very boring; the operation `yield!` will just return unit — of course, it is the way we *handle* yield that is more interesting.

```
-- Runs all of the LHS first, then the RHS.
```

```
scheduleA : {<Yield> Unit -> <Yield> Unit -> Unit}
```

```
scheduleA <yield -> m> <n> = scheduleA (m unit) n!
```

```
scheduleA <m> <yield -> n> = scheduleA m! (n unit)
```

```
scheduleA _ _ = unit
```

```
-- Lets two yields synchronise, then handles both
```



```

scheduleB : {<Yield> Unit -> <Yield> Unit -> Unit}
scheduleB <yield -> m> <yield -> n> = scheduleB (m unit) (n
    unit)
scheduleB <yield -> m> <n> = scheduleB (m unit) n!
scheduleB <m> <yield -> n> = scheduleB m! (n unit)
scheduleB _ _ = unit

```

TODO: Can maybe delete the 2nd and 3rd matches of **scheduleB** to make the point more clear?

We see two multihandlers above. Each take two **yielding** threads and schedule them, letting one run at a time. **scheduleA** runs the first thread to completion, and only then runs the second one; the first time that the second thread **yields** it is *blocked*, and can no longer execute. As such, the output of **scheduleA words! numbers!** is **one 1 two three 2 3 unit**.

scheduleB is fairer and more profound. We run **scheduleB words! number!** and receive **one 1 two 2 three 3 unit**; **scheduleB** is fair and will “match” the yields together. We step through slowly. First **words!** will print **one**, then it will **yield**. At this point — recalling that multihandlers pattern match left-to-right — the second thread, **numbers!**, is allowed to execute. In the meantime, **words!** is stuck as **<yield -> m>**; it cannot evaluate any further, it is *blocked*. Whilst **words** is blocked **numbers!** prints **1** and then **yields**. Great; now the first case matches. Both threads are resumed and the process repeats itself.

TODO: The second paragraph here is a more compelling explanation; maybe we can just get rid of all of the **scheduleA** business and /just/ have the **scheduleB** stuff? **scheduleA** is quite obvious i think whilst **B** is more subtle and compelling.

TODO: It's not true that it matches L-R as much as runs all computations L - R until they are all a command / value - fix this

1.2.2 Forking New Processes

We can make use of Frank's higher-order effects to dynamically create new threads at runtime. We strengthen the **Yield** interface by adding a new operation **fork**;

```

interface Co = fork : {[Co] Unit} -> Unit
                | yield : Unit
                | exit : Unit

```

The type of **fork** expresses that **fork** takes a suspended computation that can perform further **Co** effects, and returns unit when handled. We can now run programs that allocate new threads at runtime, such as the below

```

forker : {[Console, Co [Console]] Unit}
forker! = print "Starting! ";
          fork {print "one "; yield!; print "two "};
          fork {print "1 "; yield!; print "2 "};
          exit!

```

We can now choose a strategy for handling **fork** operations; we can either lazily run them, by finishing executing the current thread, or eagerly run them, suspending the currently executing thread and running the forked process. The handler for the former, breadth-first style of scheduling, is

```

scheduleBF : {<Co> Unit -> [Queue Proc] Unit}
scheduleBF <fork p -> k> =
    enqueue (proc {scheduleBF (<Queue> p!)});
    scheduleBF (k unit)
scheduleBF <yield -> k> =
    enqueue (proc {scheduleBF (k unit)});
    runNext!
scheduleBF <exit -> _> =
    runNext!
scheduleBF unit =
    runNext!

```

Notice the use of the **Queue** effect; we have to handle the computation **scheduleBF forker!** with a handler for **Queue** effects afterwards. Moreover, notice how concisely we can express the scheduler; this is due to the handler having access to the continuation of the caller, and treating it as a first-class object that can be stored elsewhere. We can see a diagram of how **scheduleBF** treats continuations in Figure 1.1, and a similar diagram of how the depth-first handling differs in Figure 1.2.




Figure 1.1: Breadth-First scheduling




Figure 1.2: Depth-First scheduling

Chapter 2

Formalisation of Frank

(data types)	D	(interfaces)	I
(value type variables)	X	(term variables)	x, y, z, f
(effect type variables)	E	(instance variables)	s, a, b, c
(value types)	$A, B ::= D \bar{R}$	(seeds)	$\sigma ::= \emptyset \mid E$
	$\mid \{C\} \mid X$	(abilities)	$\Sigma ::= \sigma \mid \Xi$
(computation types)	$C ::= \overline{T \rightarrow G}$	(extensions)	$\Xi ::= \mathfrak{t} \mid \Xi, I \bar{R}$
(argument types)	$T ::= \langle \Delta \rangle A$	(adaptors)	$\Theta ::= \mathfrak{t} \mid \Theta, I(S \rightarrow S')$
(return types)	$G ::= [\Sigma]A$	(adjustments)	$\Delta ::= \Theta \mid \Xi$
(type binders)	$Z ::= X \mid [E]$	(instance patterns)	$S ::= s \mid S a$
(type arguments)	$R ::= A \mid [\Sigma]$	(kind environments)	$\Phi, \Psi ::= \cdot \mid \Phi, Z$
(polytypes)	$P ::= \forall \bar{Z}. A$	(type environments)	$\Gamma ::= \cdot \mid \Gamma, x : A \mid \Gamma, f : P$
		(instance environments)	$\Omega ::= s : \Sigma \mid \Omega, a : I \bar{R}$

Figure 2.1: Types

(constructors)	k
(commands)	c
(uses)	$m ::= x \mid f \bar{R} \mid m \bar{n} \mid \uparrow(n : A)$
(constructions)	$n ::= \downarrow m \mid k \bar{n} \mid c \bar{R} \bar{n} \mid \{e\}$ $\mid \text{let } f : P = n \text{ in } n' \mid \text{letrec } \overline{f : P = e} \text{ in } n$ $\mid \langle \Theta \rangle n$
(computations)	$e ::= \overline{\bar{r} \mapsto n}$
(computation patterns)	$r ::= p \mid \langle c \bar{p} \rightarrow z \rangle \mid \langle x \rangle$
(value patterns)	$p ::= k \bar{p} \mid x$

Figure 2.2: Terms

$\Phi; \Gamma [\Sigma] \vdash m \Rightarrow A$	
$\frac{\text{T-VAR} \quad x : A \in \Gamma}{\Phi; \Gamma [\Sigma] \vdash x \Rightarrow A}$	$\frac{\text{T-POLYVAR} \quad \Phi \vdash \bar{R} \quad f : \forall \bar{Z}. A \in \Gamma}{\Phi; \Gamma [\Sigma] \vdash f \bar{R} \Rightarrow A[\bar{R}/\bar{Z}]}$
$\frac{\text{T-APP} \quad \Sigma' = \Sigma \quad (\Sigma \vdash \Delta_i \dashv \Sigma'_i)_i \quad \Phi; \Gamma [\Sigma] \vdash m \Rightarrow \{\langle \Delta \rangle A \rightarrow [\Sigma'] B\} \quad (\Phi; \Gamma [\Sigma'_i] \vdash n_i : A_i)_i}{\Phi; \Gamma [\Sigma] \vdash m \bar{n} \Rightarrow B}$	$\frac{\text{T-ASCRIBE} \quad \Phi; \Gamma [\Sigma] \vdash n : A}{\Phi; \Gamma [\Sigma] \vdash \uparrow(n : A) \Rightarrow A}$
$\Phi; \Gamma [\Sigma] \vdash n : A$	
$\frac{\text{T-SWITCH} \quad \Phi; \Gamma [\Sigma] \vdash m \Rightarrow A \quad A = B}{\Phi; \Gamma [\Sigma] \vdash \downarrow m : B}$	$\frac{\text{T-DATA} \quad k \bar{A} \in D \bar{R} \quad (\Phi; \Gamma [\Sigma] \vdash n_j : A_j)_j}{\Phi; \Gamma [\Sigma] \vdash k \bar{n} : D \bar{R}}$
$\frac{\text{T-COMMAND} \quad \Phi \vdash \bar{R} \quad c : \forall \bar{Z}. \bar{A} \rightarrow B \in \Sigma \quad (\Phi; \Gamma [\Sigma] \vdash n_j : A_j[\bar{R}/\bar{Z}])_j}{\Phi; \Gamma [\Sigma] \vdash c \bar{R} \bar{n} : B[\bar{R}/\bar{Z}]}$	$\frac{\text{T-THUNK} \quad \Phi; \Gamma \vdash e : C}{\Phi; \Gamma [\Sigma] \vdash \{e\} : \{C\}}$
$\frac{\text{T-LET} \quad P = \forall \bar{Z}. A \quad \Phi, \bar{Z}; \Gamma [\emptyset] \vdash n : A \quad \Phi; \Gamma, f : P [\Sigma] \vdash n' : B}{\Phi; \Gamma [\Sigma] \vdash \text{let } f : P = n \text{ in } n' : B}$	
$\frac{\text{T-LETREC} \quad (P_i = \forall \bar{Z}_i. \{C_i\})_i \quad (\Phi, \bar{Z}_i; \Gamma, \bar{f} : \bar{P} \vdash e_i : C)_i \quad \Phi; \Gamma, \bar{f} : \bar{P} [\Sigma] \vdash n : B}{\Phi; \Gamma [\Sigma] \vdash \text{letrec } \bar{f} : \bar{P} = e \text{ in } n : B}$	$\frac{\text{T-ADAPT} \quad \Sigma \vdash \Theta \dashv \Sigma' \quad \Phi; \Gamma [\Sigma'] \vdash n : A}{\Phi; \Gamma [\Sigma] \vdash \langle \Theta \rangle n : A}$
$\Phi; \Gamma \vdash e : C$	
$\frac{\text{T-COMP} \quad (\Phi \vdash r_{i,j} : T_j \dashv [\Sigma] \exists \Psi_{i,j}. \Gamma'_{i,j})_{i,j} \quad (\Phi, (\Psi_{i,j})_j; \Gamma, (\Gamma'_{i,j})_j [\Sigma] \vdash n_i : B)_i \quad ((r_{i,j})_i \text{ covers } T_j)_j}{\Phi; \Gamma \vdash ((r_{i,j})_j \mapsto n_i)_i : (T_j \rightarrow)_j [\Sigma] B}$	

Figure 2.3: Term Typing Rules

(uses)	$m ::= \dots \mid \lceil \mathcal{E}[c \bar{R} \bar{w}] \rceil$
(constructions)	$n ::= \dots \mid \lceil \mathcal{E}[c \bar{R} \bar{w}] \rceil$
(use values)	$u ::= x \mid f \bar{R} \mid \uparrow(v : A)$
(non-use values)	$v ::= k \bar{w} \mid \{e\}$
(construction values)	$w ::= \downarrow u \mid v$
(normal forms)	$t ::= w \mid \lceil \mathcal{E}[c \bar{R} \bar{w}] \rceil$
(evaluation frames)	$\mathcal{F} ::= [] \bar{n} \mid u (\bar{t}, [], \bar{n}) \mid \uparrow([], A)$ $\mid \downarrow[] \mid k (\bar{w}, [], \bar{n}) \mid c \bar{R} (\bar{w}, [], \bar{n})$ $\mid \mathbf{let} f : P = [] \mathbf{in} n \mid \langle \Theta \rangle []$
(evaluation contexts)	$\mathcal{E} ::= [] \mid \mathcal{F}[\mathcal{E}]$

Figure 2.4: Runtime Syntax

$\Phi; \Gamma[\Sigma] \vdash m \Rightarrow A$	$\Phi; \Gamma[\Sigma] \vdash n : A$
$\frac{\text{T-FREEZE-USE} \quad \neg(\mathcal{E} \text{ handles } c) \quad \Phi; \Gamma[\Sigma] \vdash \mathcal{E}[c \bar{R} \bar{w}] \Rightarrow A}{\Phi; \Gamma[\Sigma] \vdash \lceil \mathcal{E}[c \bar{R} \bar{w}] \rceil \Rightarrow A}$	
$\frac{\text{T-FREEZE-CONS} \quad \neg(\mathcal{E} \text{ handles } c) \quad \Phi; \Gamma[\Sigma] \vdash \mathcal{E}[c \bar{R} \bar{w}] : A}{\Phi; \Gamma[\Sigma] \vdash \lceil \mathcal{E}[c \bar{R} \bar{w}] \rceil : A}$	

Figure 2.5: Frozen Commands

$$\begin{array}{c}
\boxed{m \rightsquigarrow_{\mathbf{u}} m'} \quad \boxed{n \rightsquigarrow_{\mathbf{c}} n'} \quad \boxed{m \longrightarrow_{\mathbf{u}} m'} \quad \boxed{n \longrightarrow_{\mathbf{c}} n'} \\
\\
\text{R-HANDLE} \\
\frac{k = \min_i \{i \mid \exists \bar{\theta}. (r_{i,j} : \langle \Delta_j \rangle A_j \leftarrow t_j \neg[\Sigma] \theta_j)_j\} \quad (r_{k,j} : \langle \Delta_j \rangle A_j \leftarrow t_j \neg[\Sigma] \theta_j)_j}{\uparrow(\{((r_{i,j})_j \rightarrow n_i)_i\} : \{\langle \Delta \rangle A \rightarrow [\Sigma] B\}) \bar{t} \rightsquigarrow_{\mathbf{u}} \uparrow((\bar{\theta}(n_k) : B))} \\
\\
\begin{array}{ccc}
\text{R-ASCRIBE-USE} & \text{R-ASCRIBE-CONS} & \text{R-LET} \\
\frac{}{\uparrow(\downarrow u : A) \rightsquigarrow_{\mathbf{u}} u} & \frac{}{\downarrow \uparrow(w : A) \rightsquigarrow_{\mathbf{c}} w} & \frac{}{\mathbf{let} \ f : P = w \ \mathbf{in} \ n \rightsquigarrow_{\mathbf{c}} n[\uparrow(w : P)/f]} \\
\\
\text{R-LETREC} & \frac{}{e = \bar{r} \rightarrow n} & \text{R-ADAPT} \\
\frac{}{\mathbf{letrec} \ f : P = e \ \mathbf{in} \ n' \rightsquigarrow_{\mathbf{c}} n'[\uparrow(\{\bar{r} \rightarrow \mathbf{letrec} \ f : P = e \ \mathbf{in} \ n\} : P)/f]} & & \frac{}{\langle \Theta \rangle w \rightsquigarrow_{\mathbf{c}} w} \\
\\
\text{R-FREEZE-COMM} \\
\frac{}{c \bar{R} \bar{w} \rightsquigarrow_{\mathbf{c}} [c \bar{R} \bar{w}]} \\
\\
\begin{array}{cc}
\text{R-FREEZE-FRAME-USE} & \text{R-FREEZE-FRAME-CONS} \\
\frac{\neg(\mathcal{F}[\mathcal{E}] \text{ handles } c)}{\mathcal{F}[[\mathcal{E}[c \bar{R} \bar{w}]]] \rightsquigarrow_{\mathbf{u}} [\mathcal{F}[\mathcal{E}[c \bar{R} \bar{w}]]]} & \frac{\neg(\mathcal{F}[\mathcal{E}] \text{ handles } c)}{\mathcal{F}[[\mathcal{E}[c \bar{R} \bar{w}]]] \rightsquigarrow_{\mathbf{c}} [\mathcal{F}[\mathcal{E}[c \bar{R} \bar{w}]]]} \\
\\
\begin{array}{cccc}
\text{R-LIFT-UU} & \text{R-LIFT-UC} & \text{R-LIFT-CU} & \text{R-LIFT-CC} \\
\frac{m \rightsquigarrow_{\mathbf{u}} m'}{\mathcal{E}[m] \longrightarrow_{\mathbf{u}} \mathcal{E}[m']} & \frac{m \rightsquigarrow_{\mathbf{u}} m'}{\mathcal{E}[m] \longrightarrow_{\mathbf{c}} \mathcal{E}[m']} & \frac{n \rightsquigarrow_{\mathbf{c}} n'}{\mathcal{E}[n] \longrightarrow_{\mathbf{u}} \mathcal{E}[n']} & \frac{n \rightsquigarrow_{\mathbf{c}} n'}{\mathcal{E}[n] \longrightarrow_{\mathbf{c}} \mathcal{E}[n']}
\end{array}
\end{array}
\end{array}$$

Figure 2.6: Operational Semantics

$$\boxed{r : T \leftarrow t \neg[\Sigma] \theta}$$

$$\begin{array}{c} \text{B-VALUE} \\ \Sigma \vdash \Delta \dashv \Sigma' \\ p : A \leftarrow w \dashv \theta \\ \hline p : \langle \Delta \rangle A \leftarrow w \neg[\Sigma] \theta \end{array}$$

B-REQUEST

$$\begin{array}{c} \Sigma \vdash \Delta \dashv \Sigma' \quad \mathcal{E} \text{ poised for } c \\ \Delta = \Theta \mid \Xi \quad c : \forall \bar{Z}. \bar{B} \rightarrow B' \in \Xi \quad (p_i : B_i \leftarrow w_i \dashv \theta_i)_i \\ \hline \langle c \bar{p} \rightarrow z \rangle : \langle \Delta \rangle A \leftarrow [\mathcal{E}[c \bar{R} \bar{w}]] \neg[\Sigma] \bar{\theta}[\uparrow(\{x \mapsto \mathcal{E}[x]\} : \{B' \rightarrow [\Sigma']A\})/z] \end{array}$$

B-CATCHALL-VALUE

$$\begin{array}{c} \Sigma \vdash \Delta \dashv \Sigma' \\ \hline \langle x \rangle : \langle \Delta \rangle A \leftarrow w \neg[\Sigma] [\uparrow(\{w\} : \{[\Sigma']A\})/x] \end{array}$$

B-CATCHALL-REQUEST

$$\begin{array}{c} \Sigma \vdash \Delta \dashv \Sigma' \quad \mathcal{E} \text{ poised for } c \\ \Delta = \Theta \mid \Xi \quad c : \forall \bar{Z}. \bar{B} \rightarrow B' \in \Xi \\ \hline \langle x \rangle : \langle \Delta \rangle A \leftarrow [\mathcal{E}[c \bar{R} \bar{w}]] \neg[\Sigma] [\uparrow(\{[\mathcal{E}[c \bar{R} \bar{w}]]\} : \{[\Sigma']A\})/x] \end{array}$$

$$\boxed{p : A \leftarrow w \dashv \theta}$$

B-VAR

$$\frac{}{x : A \leftarrow w \dashv [\uparrow(w : A)/x]}$$

B-DATA

$$\frac{k \bar{A} \in D \bar{R} \quad (p_i : A_i \leftarrow w_i \dashv \theta_i)_i}{k \bar{p} : D \bar{R} \leftarrow k \bar{w} \dashv \bar{\theta}}$$

Figure 2.7: Pattern Binding

Chapter 3

Pre-emption

3.1 Motivation

One important part of our asynchronous effect handling system is the ability to interrupt arbitrary computations. This is essential for pre-emptive concurrency, which relies on being able to suspend a computation *non-cooperatively*; the computation gets suspended without being aware of its suspension.

TODO: Rewrite this - clumsy

Consider the two programs below;

```
controller : {[Stop, Go, Console] Unit}
controller! =
    stop!; print "stop "; sleep 200000; go!; controller!
```

```
runner : {[Console] Unit}
runner! = print ``1 ``; print ``2 ``; print ``3 ``;
```

We want a multihandler that uses the **stop** and **go** commands from **controller** to control the execution of **runner**. The console output of this multihandler should be then **1 stop 2 stop 3 stop**.

3.2 Interruption with Yields

We can simulate this behaviour by using the familiar **Yield** interface from Section 1.2.1.

```
runner : {[Console, Yield] Unit}
runner! = print "1 "; yield!; print "2 "; yield!; print "3 ";
    yield!
```

$$\begin{array}{c}
\text{B-CATCHALL-REQUEST} \\
\frac{\Sigma \vdash \Delta \dashv \Sigma' \quad \mathcal{E} \text{ poised for } c \quad \Delta = \Theta | \Xi \quad c : \forall \bar{Z}. \bar{B} \rightarrow B' \in \Xi}{\langle x \rangle : \langle \Delta \rangle A \leftarrow [\mathcal{E}[c \bar{R} \bar{w}]] \dashv [\Sigma] [\uparrow(\{[\mathcal{E}[c \bar{R} \bar{w}]]\} : \{[\Sigma']A\})/x]} \\
\\
\text{B-CATCHALL-REQUEST-LOOSE} \\
\frac{\Sigma \vdash \Delta \dashv \Sigma'}{\langle x \rangle : \langle \Delta \rangle A \leftarrow [\mathcal{E}[c \bar{R} \bar{w}]] \dashv [\Sigma] [\uparrow(\{[\mathcal{E}[c \bar{R} \bar{w}]]\} : \{[\Sigma']A\})/x]}
\end{array}$$

Figure 3.1: Updated B-CATCHALL-REQUEST

```

suspend : {<Yield> Unit -> <Stop, Go> Unit -> Maybe {[Console,
    Yield] Unit} -> [Console] Unit}
suspend <yield -> r> <stop -> c> _ =
    suspend unit (c unit) (just {r unit})
suspend <_> <go -> c> (just res) =
    suspend res! (c unit) nothing
suspend unit <_> _ = unit

```

Running **suspend runner! controller! nothing** then prints out **1 stop 2 stop 3** as desired. This is due to the same synchronisation behaviour that we saw in Section 1.2.1; **runner** is evaluated until it becomes a command or a value, and then **controller** is given the same treatment. Once both are a command or a value, pattern matching is done.

So far so good; this works as we hoped. However, observe that we had to change the code of the **runner** to **yield** every time it prints. This is not in the spirit of pre-emptive concurrency; we are still operating co-operatively. Threads should be unaware of the fact they are even being pre-empted. Furthermore, see that the **yield** operation adds no more information; it is just used as a placeholder operation; any operation would work. As such, we keep searching for a better solution.

TODO: Penultimate sentence could maybe go , a bit clumsy / weird

$$\begin{array}{c}
\text{B-CATCHALL-REQUEST} \\
\frac{\Sigma \vdash \Delta \dashv \Sigma' \quad \mathcal{E} \text{ poised for } c \quad \Delta = \Theta | \Xi \quad c : \forall \bar{Z}. \bar{B} \rightarrow B' \in \Xi}{\langle x \rangle : \langle \Delta \rangle A \leftarrow [\mathcal{E}[c \bar{R} \bar{w}]] \text{--}[\Sigma] [\uparrow(\{[\mathcal{E}[c \bar{R} \bar{w}]]\} : \{\Sigma' A\})/x]} \\
\\
\text{B-CATCHALL-REQUEST-LOOSE} \\
\frac{\Sigma \vdash \Delta \dashv \Sigma'}{\langle x \rangle : \langle \Delta \rangle A \leftarrow [\mathcal{E}[c \bar{R} \bar{w}]] \text{--}[\Sigma] [\uparrow(\{[\mathcal{E}[c \bar{R} \bar{w}]]\} : \{\Sigma' A\})/x]}
\end{array}$$

Figure 3.2: Updated B-CATCHALL-REQUEST

3.3 Relaxing Catches

The key to this lies in the catchall pattern, $\langle x \rangle$, and the pattern binding rules of Figure 2.7; specifically B-CATCHALL-REQUEST. We quickly go into detail on this rule now. $\langle x \rangle : \langle \Delta \rangle A$ states that $\langle x \rangle$ is a term with value type A and *adjustment* $\Delta = \Theta | \Xi$, made up of an adaptor Θ and an extension Ξ . This extension is made up of a list of interface instantiations $I \bar{R}$.

The crux is that the command c that is invoked in the frozen term $[\mathcal{E}[c \bar{R} \bar{w}]]$ must be an element of this extension Ξ ; that is, it must be handled by the current use of R-HANDLE. Refer back to the example of Section 3.2. This rule means that the catch-all pattern $\langle _ \rangle$ in the final pattern matching case of **suspend** can match against **stop** or **go**, as they are present in the extension of the second argument, but not **print** commands; although the **Console** interface is present in the ability of **controller**, it is not in the extension in **suspend**.

In the interests of pre-emption, we propose to remove this constraint from B-CATCHALL-REQUEST. The resulting rule B-CATCHALL-REQUEST-LOOSE can be seen in Figure 3.2. This lets us update the previous **suspend** code to the following, which yields the same results as last time;

```

runner : {[Console] Unit}
runner! = print "1 "; print "2 "; print "3 "

suspend : {Unit -> <Stop, Go> Unit -> Maybe {[Console] Unit}
         -> [Console] Unit}
suspend <r> <stop -> c> _ =

```

```

suspend unit (c unit) (just {r unit})
suspend <_>      <go -> c>      (just res) =
    suspend res! (c unit) nothing
suspend unit      <_>          _ = unit

```

TODO: Check that the above still works...

Now when we run `suspend runner! controller! nothing`, the `suspend` handler is able to

TODO: Talk about how this still maintains the “no-snooping” policy; we can’t inspect or access the effects that are invoked, but we know they happen.

Even with this extension,

3.4 Interrupting Arbitrary Terms

The approach of Section 3.3 can only interrupt command invocations. If `runner` were instead a sequence of pure computations¹, we would be unable to interrupt it; it does not invoke commands.

As such, we need to further change the pattern binding rules of Figure 2.7. This is to let us interrupt arbitrary computation terms. In Figure 3.3, we see an updated version of the runtime syntax; this allows for the suspension of arbitrary *uses*, being function applications and constructions.

TODO: verify that uses are “just” apps and constructions

TODO: Do we need to add interrupts to evaluation frames?

With this in mind, we now give the updated rule for the catchall pattern matching on interrupted terms (denoted $!(m)$). This can be seen in Figure 3.4. It expresses that an arbitrary suspended *use* can be matched against the computation pattern $\langle x \rangle$. The suspended computation $\{m\}$ is then bound to x , in a similar way to other B-CATCHALL rules.

Figure 3.5 shows how uses m become interrupted. This rule supplements the operational semantics of Figure 2.6. It states that at any point, a use term can reduce to the same term but suspended. At this point, the term cannot reduce any further; observe that $!()$ is not an evaluation context. The term then blocks until it is handled away. Note how similar this is to regular command invocations, which block until their continuation is invoked.

TODO: Check that just uses is enough

¹I.e. `runner! = 1 + 1; 1 + 1; 1 + 1; ...`

(uses)	$m ::= \dots \mid \lceil \mathcal{E}[c \bar{R} \bar{w}] \rceil$
(constructions)	$n ::= \dots \mid \lceil \mathcal{E}[c \bar{R} \bar{w}] \rceil$
(use values)	$u ::= x \mid f \bar{R} \mid \uparrow(v : A)$
(non-use values)	$v ::= k \bar{w} \mid \{e\}$
(construction values)	$w ::= \downarrow u \mid v$
(normal forms)	$t ::= w \mid \lceil \mathcal{E}[c \bar{R} \bar{w}] \rceil \mid !(m)$
(evaluation frames)	$\mathcal{F} ::= [] \bar{n} \mid u (\bar{t}, [], \bar{n}) \mid \uparrow([], : A)$ $\mid \downarrow[] \mid k (\bar{w}, [], \bar{n}) \mid c \bar{R} (\bar{w}, [], \bar{n})$ $\mid \text{let } f : P = [] \text{ in } n \mid \langle \Theta \rangle []$
(evaluation contexts)	$\mathcal{E} ::= [] \mid \mathcal{F}[\mathcal{E}]$

Figure 3.3: Runtime Syntax, Updated with Suspension of Uses

$$\begin{array}{c}
 \text{B-CATCHALL-INTERRUPT} \\
 \Sigma \vdash \Delta \dashv \Sigma' \\
 \hline
 \langle x \rangle : \langle \Delta \rangle A \leftarrow !(m) \dashv [\Sigma] [\uparrow(\{m\} : \{\Sigma' A\}) / x]
 \end{array}$$

Figure 3.4: Catching Interrupts rule.

TODO: Talk about how this achieves our goal.

The addition of this rule introduces non-determinism into the language; at any point, a use can either step as normal (e.g. through the R-HANDLE rule), or it can be interrupted. An interrupted term $!(m)$ can no longer reduce; it is blocked until it is resumed by the R-HANDLE rule.

TODO: Talk about non-determinism as a result of this

TODO: Maybe move non-determinism to the next section?

3.5 Interrupting In Practice

Due to the non-determinism introduced by R-INTERRUPT, this system is difficult to implement; how do we choose whether to apply a handler to its arguments or to just interrupt it?

In our implementation of this system, we instead maintain a counter which is incremented every time a handler is evaluated. When the counter reaches a certain threshold value t , we interrupt the current term m , applying the R-Interrupt rule. This converts the

R-INTERRUPT

$$\frac{}{m \rightsquigarrow_u!(m)}$$

Figure 3.5: Use interruption rule

non-deterministic system to a deterministic one; we never have any question of *what* to do.

Observe that the process of interrupting a computation is a familiar one; we stop computing and offer up the continuation to the programmer. Where is a similar control flow already around? That’s right — invocation of an effect.

TODO: Rewrite above paragraph to be less camp

As such, we can just use a normal effect to perform the interruption. Sticking with previous themes, we choose to invoke a **yield** effect when interrupting. This lets the programmer choose to handle the effect as they wish. Note that now interrupts are not restricted to the catch-all pattern $\langle m \rangle$ but are normal computation patterns.

These **yield** effects are only inserted in a computation when the **Yield** interface is present in the ability of this computation. This is important, as it lets the programmer get fine-grained control over which computations can be interrupted and which cannot. Consider the example in Section 3.1; we want the **runner** computation to be controller by the **controller**. As such, we want the **runner** to be interruptible, whilst the **controller** is not. We can reflect this by adding the **Yield** interface to the ability of the former computation.

Thus our example from before becomes;

TODO: updated example

Bibliography

Lukas Convent, Sam Lindley, Conor McBride, and Craig McLaughlin. Doo bee doo bee doo. *J. Funct. Program.*, 30:e9, 2020.

Appendix A

Remaining Formalisms

$$\boxed{\Sigma \vdash \Delta \dashv \Sigma'}$$

$$\begin{array}{c} \text{A-ADJ} \\ \Sigma \vdash \Theta \dashv \Sigma' \quad \Sigma' \vdash \Xi \dashv \Sigma'' \\ \hline \Sigma \vdash \Theta | \Xi \dashv \Sigma'' \end{array}$$

$$\boxed{\Sigma \vdash \Xi \dashv \Sigma'}$$

$$\begin{array}{c} \text{A-EXT-ID} \\ \hline \Sigma \vdash \mathbf{1} \dashv \Sigma \end{array}$$

$$\begin{array}{c} \text{A-EXT-SNOC} \\ \Sigma \vdash \Xi \dashv \Sigma' \\ \hline \Sigma \vdash \Xi, I \bar{R} \dashv \Sigma', I \bar{R} \end{array}$$

$$\boxed{\Sigma \vdash \Theta \dashv \Sigma'}$$

$$\begin{array}{c} \text{A-ADAPT-ID} \\ \hline \Sigma \vdash \mathbf{1} \dashv \Sigma \end{array}$$

$$\begin{array}{c} \text{A-ADAPT-SNOC} \\ \Sigma \vdash \Theta \dashv \Sigma' \quad \Sigma' \vdash I(S \rightarrow S') \dashv \Sigma'' \\ \hline \Sigma \vdash \Theta, I(S \rightarrow S') \dashv \Sigma'' \end{array}$$

$$\boxed{\Sigma \vdash I(S \rightarrow S') \dashv \Sigma'}$$

$$\begin{array}{c} \text{A-ADAPT-COM} \\ \Sigma \vdash S : I \dashv \Sigma'; \Omega \quad \Omega \vdash S' : I \dashv \Xi \quad \Sigma' \vdash \Xi \dashv \Sigma'' \\ \hline \Sigma \vdash I(S \rightarrow S') \dashv \Sigma'' \end{array}$$

$$\boxed{\Sigma \vdash S : I \dashv \Sigma'; \Omega}$$

$$\text{I-PAT-ID}$$

$$\Sigma \vdash s : I \dashv \Sigma; s : \Sigma$$

$$\text{I-PAT-BIND}$$

$$\Sigma \vdash S : I \dashv \Sigma'; \Omega$$

$$\Sigma, I \bar{R} \vdash S a : I \dashv \Sigma'; \Omega, a : I \bar{R}$$

$$\text{I-PAT-SKIP}$$

$$\Sigma \vdash S a : I \dashv \Sigma'; \Omega \quad I \neq I'$$

$$\Sigma, I' \bar{R} \vdash S a : I \dashv \Sigma', I' \bar{R}; \Omega$$

$$\boxed{\Omega \vdash S : I \dashv \Xi}$$

$$\begin{array}{c} \text{I-INST-ID} \\ s \in \text{dom}(\Omega) \\ \hline \Omega \vdash s : I \dashv \mathbf{1} \end{array}$$

$$\begin{array}{c} \text{I-INST-LKP} \\ a \in \text{dom}(\Omega) \quad \Omega \vdash S : I \dashv \Xi \quad \Omega(a) = I \bar{R} \\ \hline \Omega \vdash S a : I \dashv \Xi, I \bar{R} \end{array}$$

Figure A.1: Action of an Adjustment on an Ability and Auxiliary Judgements

$$\mathcal{X} ::= A \mid C \mid T \mid G \mid Z \mid R \mid P \mid \sigma \mid \Sigma \mid \Xi \mid \Theta \mid \Delta \mid \Gamma \mid \exists \Psi. \Gamma \mid \Omega$$

$\Phi \vdash \mathcal{X}$

$\frac{}{\Phi, X \vdash X}$ <p>WF-VAL</p>	$\frac{}{\Phi, [E] \vdash E}$ <p>WF-EFF</p>	$\frac{\Phi, \bar{Z} \vdash A}{\Phi \vdash \forall \bar{Z}. A}$ <p>WF-POLY</p>
$\frac{(\Phi \vdash R)_i}{\Phi \vdash D \bar{R}}$ <p>WF-DATA</p>	$\frac{\Phi \vdash C}{\Phi \vdash \{C\}}$ <p>WF-THUNK</p>	$\frac{(\Phi \vdash T)_i \quad \Phi \vdash G}{\Phi \vdash \bar{T} \rightarrow G}$ <p>WF-COMP</p>
$\frac{\Phi \vdash \Delta \quad \Phi \vdash A}{\Phi \vdash \langle \Delta \rangle A}$ <p>WF-ARG</p>		
$\frac{\Phi \vdash \Sigma \quad \Phi \vdash A}{\Phi \vdash [\Sigma] A}$ <p>WF-RET</p>	$\frac{\Phi \vdash \Sigma}{\Phi \vdash [\Sigma]}$ <p>WF-ABILITY</p>	$\frac{}{\Phi \vdash \emptyset}$ <p>WF-PURE</p>
$\frac{}{\Phi \vdash \mathbf{1}}$ <p>WF-ID</p>		$\frac{\Phi \vdash \Xi \quad (\Phi \vdash R)_i}{\Phi \vdash \Xi, I \bar{R}}$ <p>WF-EXT</p>
$\frac{\Phi \vdash \Theta}{\Phi \vdash \Theta, I (S \rightarrow S')}$ <p>WF-ADAPT</p>		
$\frac{}{\Phi \vdash \cdot}$ <p>WF-EMPTY</p>	$\frac{\Phi \vdash \Gamma \quad \Phi \vdash A}{\Phi \vdash \Gamma, x : A}$ <p>WF-MONO</p>	$\frac{\Phi \vdash \Gamma \quad \Phi \vdash P}{\Phi \vdash \Gamma, f : P}$ <p>WF-POLY</p>
$\frac{\Phi, \Psi \vdash \Gamma}{\Phi \vdash \exists \Psi. \Gamma}$ <p>WF-EXISTENTIAL</p>		$\frac{\Phi \vdash \Omega \quad (\Phi \vdash R)_i}{\Phi \vdash \Omega, x : I \bar{R}}$ <p>WF-INTERFACE</p>

Figure A.2: Well-Formedness Rules

$$\boxed{\Phi \vdash p : A \dashv \Gamma}$$

$$\begin{array}{c}
\text{P-VAR} \\
\hline
\Phi \vdash x : A \dashv x : A
\end{array}
\qquad
\begin{array}{c}
\text{P-DATA} \\
\frac{k \bar{A} \in D \bar{R} \quad (\Phi \vdash p_i : A_i \dashv \Gamma)_i}{\Phi \vdash k \bar{p} : D \bar{R} \dashv \bar{\Gamma}}
\end{array}$$

$$\boxed{\Phi \vdash r : T \dashv [\Sigma] \exists \Psi. \Gamma}$$

$$\begin{array}{c}
\text{P-VALUE} \\
\frac{\Sigma \vdash \Delta \dashv \Sigma' \quad \Phi \vdash p : A \dashv \Gamma}{\Phi \vdash p : \langle \Delta \rangle A \dashv [\Sigma] \Gamma}
\end{array}
\qquad
\begin{array}{c}
\text{P-CATCHALL} \\
\frac{\Sigma \vdash \Delta \dashv \Sigma'}{\Phi \vdash \langle x \rangle : \langle \Delta \rangle A \dashv [\Sigma] x : \{[\Sigma'] A\}}
\end{array}$$

$$\begin{array}{c}
\text{P-COMMAND} \\
\frac{\Sigma \vdash \Delta \dashv \Sigma' \quad \Delta = \Theta \mid \Xi \quad c : \forall \bar{Z}. \bar{A} \rightarrow \bar{B} \in \Xi \quad (\Phi, \bar{Z} \vdash p_i : A_i \dashv \Gamma_i)_i}{\Phi \vdash \langle c \bar{p} \rightarrow z \rangle : \langle \Delta \rangle B' \dashv [\Sigma] \exists \bar{Z}. \bar{\Gamma}, z : \{\langle \mathbf{1} \mid \mathbf{1} \rangle B \rightarrow [\Sigma'] B'\}}
\end{array}$$

Figure A.3: Pattern Matching Typing Rules